

Are Facebook likes, profits and movie reviews accurate predictors for IMDB score?

kvk229 & gsc326

11/22/2016

GAM model general additive model

log linear regression

```
# setting the working directory
setwd('/Users/kkiran/Desktop/fall_2016/fds/project/MovieScorePredictor/data')
# 'C:/Users/gogs/Documents/GitHub/MovieScorePredictor/Data'
# setwd('C:/Users/gogs/Documents/GitHub/MovieScorePredictor/Data')
movieData = read.xls("movie_data.xls")
#head(movieData)
```

Loaded the data into a data frame 'movieData'

```
# identifying top 10 genres out of all the 26 genres to make the work more focussed
```

```
# gernes present in the data:
```

```
# 'Sci-Fi', 'Crime', 'Romance', 'Animation', 'Music', 'Comedy', 'War', 'genres', 'Horror', 'Film-Noir', 'Adventure'
```

```
movieCount <- c()
```

```
for(i in 38:64)
```

```
{
```

```
  movieCount[i - 37] = sum(movieData[,i]);
```

```
}
```

```
movieCount
```

```
## [1] 616 889 1107 242 214 1872 213 0 565 6 923 3 2 1411
```

```
## [15] 97 500 5 2594 1153 121 132 207 546 610 1 182 293
```

```
genreNames <- as.vector(colnames(movieData)[38:64])
```

```
genreNames
```

```
## [1] "Sci-Fi"
```

```
"Crime"
```

```
"Romance"
```

```
"Animation"
```

```
"Music"
```

```
## [6] "Comedy"
```

```
"War"
```

```
"genres.1"
```

```
"Horror"
```

```
"Film.Noir"
```

```
## [11] "Adventure"
```

```
"News"
```

```
"Reality.TV"
```

```
"Thriller"
```

```
"Western"
```

```
## [16] "Mystery"
```

```
"Short"
```

```
"Drama"
```

```
"Action"
```

```
"Documentary"
```

```
## [21] "Musical"
```

```
"History"
```

```
"Family"
```

```
"Fantasy"
```


```
"Game.Show"
```

```
## [26] "Sport"
```

```
"Biography"
```

```
genreNames <- as.vector(genreNames)


genreMovieCount <- data.frame(genreNames, movieCount)
plot(genreMovieCount$genreNames, genreMovieCount$movieCount, main="Genre Distribution", xlab="Genre ", ylab="Movies Made")
```



We can see that not all the genres have a considerable number of movies made in them, so we decided to extract the top 11 genres that have the most number of movies made in those particular genres.

```
genreMovieCountSorted <- genreMovieCount[order(-movieCount),]
genreMovieCountSorted <- genreMovieCountSorted[c(1:10),]
plot(genreMovieCountSorted$genreNames, genreMovieCountSorted$movieCount, main="Filtered Genre Distribution", xlab="Genre ", ylab="Movies Made")
```

Filtered Genre Distribution



Now we have to remove the data for all the other genres from the data set, also we can delete the last column from the data set because it is repeated

```
movieData <- movieData[,-65]
columnNames <- colnames(movieData)
columnNames <- columnNames[1:37]

selectedNames <- genreMovieCountSorted$genreNames

columnNames <- as.vector(columnNames)
selectedNames <- as.vector(selectedNames)

names <- c(columnNames, selectedNames)
names
```

```
## [1] "movie_title" "actor_1_facebook_likes"
## [3] "actor_2_facebook_likes" "actor_3_facebook_likes"
## [5] "director_facebook_likes" "movie_facebook_likes"
## [7] "cast_total_facebook_likes" "director_name"
## [9] "actor_1_name" "actor_2_name"
## [11] "actor_3_name" "gross"
## [13] "budget" "imdb_score"
## [15] "num_critic_for_reviews" "num_user_for_reviews"
## [17] "tomatoUserRating" "tomatoRating"
## [19] "tomatoReviews" "tomatoFresh"
## [21] "tomatoRotten" "tomatoUserMeter"
## [23] "tomatoUserReviews" "num_voted_users"
## [25] "imdbVotes" "Metascore"
## [27] "genres" "facenumber_in_poster"
```

```
## [29] "plot_keywords"      "movie_imdb_link"
## [31] "language"           "country"
## [33] "content_rating"     "title_year"
## [35] "aspect_ratio"       "color"
## [37] "duration"           "Drama"
## [39] "Comedy"              "Thriller"
## [41] "Action"              "Romance"
## [43] "Adventure"           "Crime"
## [45] "Sci.Fi"              "Fantasy"
## [47] "Horror"
```

```
movieData1 <- subset(movieData, select = names)
# for(i in 38:47)
# {
#   print(c(colnames(movieData1[i]), sum(movieData1[,i])));
# }
```

Now movieData1 has the data about the genres that we are interested only. The next step is to clean the data by removing rows that dont have a considerable amount of data. If the facebook likes are missing, although we can get those details from other row, we are not proceeding so because the number of facebook likes are always changing and fetching data from other rows might not be a very good estimate for plugging in missing likes data.

Note: The facebook data is present in the data itself, we didnt have to fetch the data manually.

```
mean(is.na(movieData))
```

```
## [1] 0.006379511
```

```
paste("only", mean(is.na(movieData)), " (mean) amount of data is null, so we can safely remove NAs")
```

```
## [1] "only 0.00637951120364862 (mean) amount of data is null, so we can safely remove NAs"
```

```
row.has.na <- apply(movieData1, 1, function(x){any(is.na(x))})
numberOfNAs <- sum(row.has.na)

print (c("can remove ", numberOfNAs , "null rows from the table"))
```

```
## [1] "can remove "      "1242"
## [3] "null rows from the table"
```

```
#removing the nulls
movieData1 <- na.omit(movieData1)
```

```
NACounter <- 0
indicesToRemove <- c()
index <- 1
# this is the working version
for (i in 1 : nrow(movieData1)) {
  if (any(movieData1[i,] == "N/A")) {
    # print (c(i, "yes" , movieData1[i,]))
  }
}
```

```

    indicesToRemove[index] = i;
    index <- index + 1
    NAcouter <- NAcouter + 1
  }
  # print ("no")
}

print(length(indicesToRemove))

```

```
## [1] 567
```

```

# print(indicesToRemove)
print(c("total number of nulls", NAcouter))

```

```
## [1] "total number of nulls" "567"
```

```

#removing the 'NAcouter' number of rows that have NA in them
movieData2 <- movieData1[-indicesToRemove,]
# Now, movieData2 has no NA in any of the rows.

```

```

movieData2[movieData2==""] <- NA
row.has.na <- apply(movieData2, 1, function(x){any(is.na(x))})
numberOfNAs <- sum(row.has.na)
paste("There are ",numberOfNAs, " rows with empty cells, so we are removing them")

```

```
## [1] "There are 236 rows with empty cells, so we are removing them"
```

```

# removing the empty rows, (second round of filtering)
movieData2 <- na.omit(movieData2)

```

```

#counting the profits of a movie by subtracting the budget from the gross
movieData2$profits <- movieData2$gross - movieData2$budget
movieData3 <- movieData2[,c(c(1:13),48, c(14:47))]

```

```

# movieData3 is the final cleaned data that also has a column showing the profits made by the movie
#str(movieData3)
stat <- nearZeroVar(movieData3, saveMetrics = T)
class(stat$zeroVar)

```

```
## [1] "logical"
```

```
varDF <- cbind.data.frame(colnames(movieData3),stat$zeroVar)
```

```

#converting logical to binary
cols <- sapply(varDF, is.logical)
varDF[,cols] <- lapply(varDF[,cols], as.numeric)

```

```


## Warning in `[<-data.frame`(`*tmp*`, , cols, value = list(0, 0, 0, 0, 0, :
## provided 48 variables to replace 1 variables

```

```
# for ( i in 1:nrow(varDF)) {
#   print(varDF[i,2])
# }
```

```
plot(varDF$`colnames(movieData3)`, varDF$`stat$zeroVar`, xlab = "predictors", ylab = "Zero - variance"
```

zero variance scores of different possible predictors



```
paste("we can see that all the columns have a zero score for the zero-variance score showing that all t
```

```
## [1] "we can see that all the columns have a zero score for the zero-variance score showing that all t
```

```
# We can see that none of the variables have zero variance, so we can consider all the variables for stu
```

```
#converting factor to int array
movieData3$tomatoUserRating <- as.numeric(as.character(movieData3$tomatoUserRating))
movieData3$tomatoRating <- as.numeric(as.character(movieData3$tomatoRating))
movieData3$tomatoReviews <- as.numeric(as.character(movieData3$tomatoReviews))
movieData3$tomatoFresh <- as.numeric(as.character(movieData3$tomatoFresh))
movieData3$tomatoRotten <- as.numeric(as.character(movieData3$tomatoRotten))
movieData3$tomatoUserMeter <- as.numeric(as.character(movieData3$tomatoUserMeter))
movieData3$tomatoUserReviews <- as.numeric(as.character(movieData3$tomatoUserReviews))
movieData3$imdbVotes <- as.numeric(as.character(movieData3$imdbVotes))
movieData3$Metascore <- as.numeric(as.character(movieData3$Metascore))
```

```
# checking the correlation
```

```
# str(movieData3)
```

```
cor(movieData3$imdb_score, movieData3[,c(7,12:14,16:26,38)], use = "pairwise.complete.obs")
```

```
##      cast_total_facebook_likes      gross      budget      profits
## [1,]                0.0980045 0.2074095 0.02979009 0.02925642
##      num_critic_for_reviews num_user_for_reviews tomatoUserRating
## [1,]                0.3549605                0.3305453                0.8103758
##      tomatoRating tomatoReviews tomatoFresh tomatoRotten tomatoUserMeter
## [1,]                0.813982                0.2975506                0.5848666 -0.4436326                0.8440624
##      tomatoUserReviews num_voted_users imdbVotes      duration
## [1,]                0.08724109                0.4919135 0.4907566 0.3701765
```

```
colnames(movieData3)[c(7,12:14,16:26,38)]
```

```
## [1] "cast_total_facebook_likes" "gross"
## [3] "budget"                    "profits"
## [5] "num_critic_for_reviews"    "num_user_for_reviews"
## [7] "tomatoUserRating"          "tomatoRating"
## [9] "tomatoReviews"             "tomatoFresh"
## [11] "tomatoRotten"              "tomatoUserMeter"
## [13] "tomatoUserReviews"         "num_voted_users"
## [15] "imdbVotes"                 "duration"
```

```
cNames <- paste("movieData3$",colnames(movieData3)[c(7,12:14,16:26,38)] , sep = "")
# cNames
#formula contains all the columns that we want to include in the model
formula <- as.formula(paste("y ~ ", paste(cNames, collapse= "+")))
# formula
```

```
#choice of linear regression vs logistic regression
```

```
# Linear regression: When the outcome(dependent variable) is continuous, i.e. infinite number of possib
```

```
# Logistic regression: When the outcome(dependent variable) has a limited set of values.
```

```
#because we are trying to predict the IMDB score of a movie and theoritically the score can have an inf
```

```
#linear regression
```

```
lmfit1.movieData <- lm(movieData3$imdb_score ~ movieData3$cast_total_facebook_likes + movieData3$gross +
```

```
# summary(lmfit1.movieData)
```

```
vif(lmfit1.movieData)
```

```
## movieData3$cast_total_facebook_likes      movieData3$gross
##                                1.104381                1.780129
##      movieData3$budget      movieData3$num_critic_for_reviews
##                                1.022601                4.204645
##      movieData3$num_user_for_reviews      movieData3$tomatoUserRating
##                                3.200089                6.839647
##      movieData3$tomatoRating      movieData3$tomatoReviews
##                                5.104971                7.447773
##      movieData3$tomatoFresh      movieData3$tomatoUserMeter
##                                11.679957                7.997424
##      movieData3$tomatoUserReviews      movieData3$num_voted_users
##                                1.165426                364.082025
##      movieData3$imdbVotes      movieData3$duration
##                                370.657616                1.286969
```

Looking at the Variance Inflation Factor of the fitted model, we can see that imdbVotes and num_voted_users are two largest VIFs. so we try to eliminate them and make the model again.

```
# coefficients of multi variate linear regression  
paste("coefficients of fitted line by linear regression")
```

```
## [1] "coefficients of fitted line by linear regression"
```

```
# summary(lmfit1.movieData)$coefficients
```

```
lmfit2.movieData <- lm(movieData3$imdb_score ~ movieData3$cast_total_facebook_likes + movieData3$gross  
  
# summary(lmfit2.movieData)  
  
lmfit2.movieData
```

```
##  
## Call:  
## lm(formula = movieData3$imdb_score ~ movieData3$cast_total_facebook_likes +  
##     movieData3$gross + movieData3$budget + movieData3$num_critic_for_reviews +  
##     movieData3$num_user_for_reviews + movieData3$tomatoUserRating +  
##     +movieData3$tomatoRating + movieData3$tomatoReviews + movieData3$tomatoFresh +  
##     movieData3$tomatoUserMeter + movieData3$tomatoUserReviews +  
##     movieData3$duration, data = movieData3)  
##  
## Coefficients:  
##                (Intercept)  
##                1.103e+00  
## movieData3$cast_total_facebook_likes  
##                1.633e-06  
##                movieData3$gross  
##                -1.043e-09  
##                movieData3$budget  
##                7.982e-11  
## movieData3$num_critic_for_reviews  
##                2.418e-04  
## movieData3$num_user_for_reviews  
##                1.256e-04  
## movieData3$tomatoUserRating  
##                5.397e-01  
## movieData3$tomatoRating  
##                3.519e-01  
## movieData3$tomatoReviews  
##                3.114e-03  
## movieData3$tomatoFresh  
##                -4.768e-03  
## movieData3$tomatoUserMeter  
##                1.857e-02  
## movieData3$tomatoUserReviews  
##                1.389e-08  
## movieData3$duration  
##                2.786e-03
```



```
vif(lmfit2.movieData)
```

```
## movieData3$cast_total_facebook_likes      movieData3$gross
##                1.087276                1.583824
##                movieData3$budget      movieData3$num_critic_for_reviews
##                1.019983                4.043964
##                movieData3$num_user_for_reviews      movieData3$tomatoUserRating
##                2.072406                6.694771
##                movieData3$tomatoRating      movieData3$tomatoReviews
##                5.104545                7.339284
##                movieData3$tomatoFresh      movieData3$tomatoUserMeter
##                11.627723                7.993222
##                movieData3$tomatoUserReviews      movieData3$duration
##                1.161993                1.286622
```

```
# we can see that the std error for the parameter estimates gets smaller.
```

```
#calculating the MSFE and MAD for the predicted values
```

```
predictedScore <- predict(lmfit2.movieData)
# predictedScore
RSFE_v <- movieData3$imdb_score - predictedScore
# RSFE_v
RSFE <- sum(RSFE_v)
# RSFE
absRSFE <- abs(RSFE)
# absRSFE
length(RSFE_v)
```

```
## [1] 2998
```

```
MSFE <- absRSFE / length(RSFE_v)
# MSFE
mean(lmfit2.movieData$residuals^2)
```

```
## [1] 0.1941867
```

```
# calculating mad
```

```
Madoriginal <- mad(movieData3$imdb_score,center = median(movieData3$imdb_score),constant = 1)
MadRegression <- mad(predictedScore, center = median(predictedScore), constant = 1)
```

```
# Analysing the linear regression model
```

```
# log.movieData <- log(movieData3[, c(7,12:26,38)])
trans = preprocess(x = movieData3[,c(7,12:14,16:26,38)], method=c("BoxCox", "center", "scale", "pca"))
trans
```

```
## Created from 2998 samples and 16 variables
##
## Pre-processing:
##   - Box-Cox transformation (13)
##   - centered (16)
##   - ignored (0)
```

```

## - principal component signal extraction (16)
## - scaled (16)
##
## Lambda estimates for Box-Cox transformation:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -0.8000  0.1000  0.2000  0.3308  0.5000  1.3000
##
## PCA needed 9 components to capture 95 percent of the variance

trans$pcaComp

## NULL

pca.movieData <- prcomp(x = movieData3[,c(7,12:14,16:26,38)], center = TRUE, scale. = TRUE)
print(pca.movieData)

## Standard deviations:
## [1] 2.432503e+00 1.564950e+00 1.401758e+00 1.154068e+00 9.799094e-01
## [6] 9.184594e-01 8.827092e-01 7.648408e-01 7.043480e-01 5.112467e-01
## [11] 4.578182e-01 3.645493e-01 2.603347e-01 3.693681e-02 1.698190e-15
## [16] 9.873104e-16
##
## Rotation:
##
##               PC1          PC2          PC3          PC4
## cast_total_facebook_likes 0.10874858 -0.159739097 -0.03150854 -0.07741693
## gross                     0.26316009 -0.208146502 -0.08207119  0.18149678
## budget                    0.02117848 -0.124725036  0.68823554  0.09564829
## profits                   0.05347272  0.063947013 -0.70036681 -0.04286223
## num_critic_for_reviews    0.31157844 -0.195203809  0.03556094 -0.34827959
## num_user_for_reviews      0.30429167 -0.199198393 -0.03953276  0.26783237
## tomatoUserRating          0.27568875  0.348381799  0.07634384 -0.09588632
## tomatoRating              0.28271713  0.375475248  0.05987576 -0.10215443
## tomatoReviews             0.28800338 -0.262819696  0.01164888 -0.44391120
## tomatoFresh               0.34435948  0.048346260  0.02976664 -0.34734537
## tomatoRotten              -0.06553253 -0.535196344 -0.02864047 -0.19999098
## tomatoUserMeter           0.26261601  0.413375677  0.06884329  0.02569731
## tomatoUserReviews         0.09772351 -0.126815013 -0.06026578  0.44463194
## num_voted_users           0.35162157 -0.114814179 -0.03940306  0.27747165
## imdbVotes                 0.35267478 -0.115679235 -0.03700495  0.27411764
## duration                  0.19402743 -0.002161922  0.05209069  0.17095599
##
##               PC5          PC6          PC7
## cast_total_facebook_likes 0.7442970309 0.57131652 0.199802020
## gross                     0.0371504409 0.07906111 -0.152176368
## budget                    -0.0007362571 0.05667193 -0.053212424
## profits                   0.0112150129 -0.03343718 0.009386848
## num_critic_for_reviews    -0.1138773738 -0.02993015 -0.049267286
## num_user_for_reviews      -0.0526316998 -0.15005822 -0.128820954
## tomatoUserRating          0.1329041063 -0.02261445 0.033480422
## tomatoRating              -0.1115986321 0.09782440 0.079849278
## tomatoReviews             -0.1770780212 0.01277372 0.140547048
## tomatoFresh               -0.2268091964 0.12742434 0.072353255
## tomatoRotten              0.0650392947 -0.18694074 0.125051821
## tomatoUserMeter           0.1129269225 -0.01149642 0.046752639

```

## tomatoUserReviews	-0.4651400109	0.45541260	0.538633766
## num_voted_users	0.0748453400	-0.04572894	-0.289635505
## imdbVotes	0.0739755760	-0.04695573	-0.290051685
## duration	0.2672984666	-0.60526603	0.637593608
##	PC8	PC9	PC10
## cast_total_facebook_likes	0.17506954	0.020426054	-0.05914578
## gross	-0.49744154	-0.713144917	-0.12911647
## budget	-0.07563258	-0.126002502	-0.02731393
## profits	-0.06604902	-0.077397867	-0.00958438
## num_critic_for_reviews	0.13444239	0.001216694	0.31909171
## num_user_for_reviews	0.28194783	0.201509011	-0.72552433
## tomatoUserRating	-0.47326075	0.279209041	0.10658684
## tomatoRating	0.22778587	-0.159635157	-0.12269418
## tomatoReviews	-0.05719630	0.051030136	-0.09529204
## tomatoFresh	0.17192690	-0.149931431	-0.03681209
## tomatoRotten	-0.38133069	0.334541978	-0.10487794
## tomatoUserMeter	-0.33876910	0.253709205	-0.24588388
## tomatoUserReviews	-0.10465477	0.163945599	0.11853168
## num_voted_users	0.08615309	0.164713071	0.32513376
## imdbVotes	0.08207903	0.162650775	0.32435965
## duration	0.14236754	-0.207984553	0.13231874
##	PC11	PC12	PC13
## cast_total_facebook_likes	0.026671700	0.0054614142	-0.007909208
## gross	0.057016165	-0.0444305061	0.009528170
## budget	-0.005351173	-0.0004943584	-0.001489339
## profits	0.021365572	-0.0120598955	0.004156011
## num_critic_for_reviews	0.714951697	-0.1901441546	0.242204070
## num_user_for_reviews	0.267208970	0.0549783993	-0.177867179
## tomatoUserRating	0.223888668	0.0837567248	-0.630769826
## tomatoRating	-0.232555094	-0.7595995477	-0.114768959
## tomatoReviews	-0.325220711	0.1710452076	-0.024027007
## tomatoFresh	-0.223312546	0.4218467584	-0.039739387
## tomatoRotten	-0.197653057	-0.3955499700	0.023535302
## tomatoUserMeter	-0.017069174	0.0747094918	0.703002764
## tomatoUserReviews	0.082555510	-0.0190121060	0.009940902
## num_voted_users	-0.238176326	-0.0055349746	0.024580202
## imdbVotes	-0.223861943	-0.0111021480	0.028969046
## duration	-0.011638249	0.0699895743	0.004428365
##	PC14	PC15	PC16
## cast_total_facebook_likes	-0.0009811830	3.250883e-18	6.205474e-17
## gross	0.0030325036	-2.761084e-02	1.953830e-01
## budget	0.0013590672	9.620970e-02	-6.808101e-01
## profits	-0.0004809376	9.777874e-02	-6.919131e-01
## num_critic_for_reviews	0.0095219764	-1.222113e-15	-3.261280e-16
## num_user_for_reviews	0.0010740188	-4.961309e-16	1.387779e-16
## tomatoUserRating	0.0011539765	-4.644722e-16	5.828671e-16
## tomatoRating	0.0002088994	1.474515e-16	-3.035766e-16
## tomatoReviews	-0.0023931164	6.634452e-01	9.375575e-02
## tomatoFresh	-0.0031922237	-6.276595e-01	-8.869865e-02
## tomatoRotten	0.0010873997	-3.824979e-01	-5.405326e-02
## tomatoUserMeter	0.0023808116	2.515349e-16	-7.979728e-17
## tomatoUserReviews	-0.0003350411	1.006140e-16	1.327063e-16
## num_voted_users	0.7038662332	1.040834e-17	-2.151057e-16
## imdbVotes	-0.7102420018	4.770490e-16	3.330669e-16

```
## duration -0.0005571525 1.305379e-16 -5.551115e-17
```

```
colnames(pca.movieData$rotation)
```

```
## [1] "PC1" "PC2" "PC3" "PC4" "PC5" "PC6" "PC7" "PC8" "PC9" "PC10"  
## [11] "PC11" "PC12" "PC13" "PC14" "PC15" "PC16"
```

```
# summary(pca.movieData)
```

```
pca.movieData
```

```
## Standard deviations:
```

```
## [1] 2.432503e+00 1.564950e+00 1.401758e+00 1.154068e+00 9.799094e-01  
## [6] 9.184594e-01 8.827092e-01 7.648408e-01 7.043480e-01 5.112467e-01  
## [11] 4.578182e-01 3.645493e-01 2.603347e-01 3.693681e-02 1.698190e-15  
## [16] 9.873104e-16  
##
```


```
## Rotation:
```

```
##          PC1          PC2          PC3          PC4  
## cast_total_facebook_likes 0.10874858 -0.159739097 -0.03150854 -0.07741693  
## gross 0.26316009 -0.208146502 -0.08207119 0.18149678  
## budget 0.02117848 -0.124725036 0.68823554 0.09564829  
## profits 0.05347272 0.063947013 -0.70036681 -0.04286223  
## num_critic_for_reviews 0.31157844 -0.195203809 0.03556094 -0.34827959  
## num_user_for_reviews 0.30429167 -0.199198393 -0.03953276 0.26783237  
## tomatoUserRating 0.27568875 0.348381799 0.07634384 -0.09588632  
## tomatoRating 0.28271713 0.375475248 0.05987576 -0.10215443  
## tomatoReviews 0.28800338 -0.262819696 0.01164888 -0.44391120  
## tomatoFresh 0.34435948 0.048346260 0.02976664 -0.34734537  
## tomatoRotten -0.06553253 -0.535196344 -0.02864047 -0.19999098  
## tomatoUserMeter 0.26261601 0.413375677 0.06884329 0.02569731  
## tomatoUserReviews 0.09772351 -0.126815013 -0.06026578 0.44463194  
## num_voted_users 0.35162157 -0.114814179 -0.03940306 0.27747165  
## imdbVotes 0.35267478 -0.115679235 -0.03700495 0.27411764  
## duration 0.19402743 -0.002161922 0.05209069 0.17095599  
##          PC5          PC6          PC7  
## cast_total_facebook_likes 0.7442970309 0.57131652 0.199802020  
## gross 0.0371504409 0.07906111 -0.152176368  
## budget -0.0007362571 0.05667193 -0.053212424  
## profits 0.0112150129 -0.03343718 0.009386848  
## num_critic_for_reviews -0.1138773738 -0.02993015 -0.049267286  
## num_user_for_reviews -0.0526316998 -0.15005822 -0.128820954  
## tomatoUserRating 0.1329041063 -0.02261445 0.033480422  
## tomatoRating -0.1115986321 0.09782440 0.079849278  
## tomatoReviews -0.1770780212 0.01277372 0.140547048  
## tomatoFresh -0.2268091964 0.12742434 0.072353255  
## tomatoRotten 0.0650392947 -0.18694074 0.125051821  
## tomatoUserMeter 0.1129269225 -0.01149642 0.046752639  
## tomatoUserReviews -0.4651400109 0.45541260 0.538633766  
## num_voted_users 0.0748453400 -0.04572894 -0.289635505  
## imdbVotes 0.0739755760 -0.04695573 -0.290051685  
## duration 0.2672984666 -0.60526603 0.637593608  
##          PC8          PC9          PC10  
## cast_total_facebook_likes 0.17506954 0.020426054 -0.05914578
```

## gross	-0.49744154	-0.713144917	-0.12911647
## budget	-0.07563258	-0.126002502	-0.02731393
## profits	-0.06604902	-0.077397867	-0.00958438
## num_critic_for_reviews	0.13444239	0.001216694	0.31909171
## num_user_for_reviews	0.28194783	0.201509011	-0.72552433
## tomatoUserRating	-0.47326075	0.279209041	0.10658684
## tomatoRating	0.22778587	-0.159635157	-0.12269418
## tomatoReviews	-0.05719630	0.051030136	-0.09529204
## tomatoFresh	0.17192690	-0.149931431	-0.03681209
## tomatoRotten	-0.38133069	0.334541978	-0.10487794
## tomatoUserMeter	-0.33876910	0.253709205	-0.24588388
## tomatoUserReviews	-0.10465477	0.163945599	0.11853168
## num_voted_users	0.08615309	0.164713071	0.32513376
## imdbVotes	0.08207903	0.162650775	0.32435965
## duration	0.14236754	-0.207984553	0.13231874
##	PC11	PC12	PC13
## cast_total_facebook_likes	0.026671700	0.0054614142	-0.007909208
## gross	0.057016165	-0.0444305061	0.009528170
## budget	-0.005351173	-0.0004943584	-0.001489339
## profits	0.021365572	-0.0120598955	0.004156011
## num_critic_for_reviews	0.714951697	-0.1901441546	0.242204070
## num_user_for_reviews	0.267208970	0.0549783993	-0.177867179
## tomatoUserRating	0.223888668	0.0837567248	-0.630769826
## tomatoRating	-0.232555094	-0.7595995477	-0.114768959
## tomatoReviews	-0.325220711	0.1710452076	-0.024027007
## tomatoFresh	-0.223312546	0.4218467584	-0.039739387
## tomatoRotten	-0.197653057	-0.3955499700	0.023535302
## tomatoUserMeter	-0.017069174	0.0747094918	0.703002764
## tomatoUserReviews	0.082555510	-0.0190121060	0.009940902
## num_voted_users	-0.238176326	-0.0055349746	0.024580202
## imdbVotes	-0.223861943	-0.0111021480	0.028969046
## duration	-0.011638249	0.0699895743	0.004428365
##	PC14	PC15	PC16
## cast_total_facebook_likes	-0.0009811830	3.250883e-18	6.205474e-17
## gross	0.0030325036	-2.761084e-02	1.953830e-01
## budget	0.0013590672	9.620970e-02	-6.808101e-01
## profits	-0.0004809376	9.777874e-02	-6.919131e-01
## num_critic_for_reviews	0.0095219764	-1.222113e-15	-3.261280e-16
## num_user_for_reviews	0.0010740188	-4.961309e-16	1.387779e-16
## tomatoUserRating	0.0011539765	-4.644722e-16	5.828671e-16
## tomatoRating	0.0002088994	1.474515e-16	-3.035766e-16
## tomatoReviews	-0.0023931164	6.634452e-01	9.375575e-02
## tomatoFresh	-0.0031922237	-6.276595e-01	-8.869865e-02
## tomatoRotten	0.0010873997	-3.824979e-01	-5.405326e-02
## tomatoUserMeter	0.0023808116	2.515349e-16	-7.979728e-17
## tomatoUserReviews	-0.0003350411	1.006140e-16	1.327063e-16
## num_voted_users	0.7038662332	1.040834e-17	-2.151057e-16
## imdbVotes	-0.7102420018	4.770490e-16	3.330669e-16
## duration	-0.0005571525	1.305379e-16	-5.551115e-17

```
plot(pca.movieData, type = "l", main = "PCA Movie Data")
```

PCA Movie Data



```
wssplot <- function(data, nc=15, seed=1234) {
  wss <- (nrow(data)-1)*sum(apply(data,2,var))
  bss <- (nrow(data)-1)*sum(apply(data,2,var))
  for (i in 2:nc) {
    set.seed(seed)
    wss[i] <- sum(kmeans(data[,c(7,12:26,38)], centers=i)$withinss)
    bss[i] <- sum(kmeans(data[,c(7,12:26,38)], centers=i)$betweenss)
  }


  plot(1:nc, wss, type="b", xlab="Number of Clusters",
       ylab="Within groups sum of squares")
  plot(1:nc, bss, type="b", xlab="Number of Clusters",
       ylab="Between groups sum of squares")
}


wssplot(movieData3, nc=35)
```

```
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
```

```
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning in FUN(newX[, i], ...): NAs introduced by coercion
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
```

```
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
## Warning: did not converge in 10 iterations
```





```
paste("from both of the graphs above, we can see that 20 is the optimum number of clusters because the l
```

```
## [1] "from both of the graphs above, we can see that 20 is the optimum number of clusters because the
```

```
paste("The within groups sum of squares was decreasing untill 20 clusters and the between groups sum of
```

```
## [1] "The within groups sum of squares was decreasing untill 20 clusters and the between groups sum o
```

```
set.seed(101)
sample = sample.split(movieData3, SplitRatio = .75)
train = subset(movieData3, sample == TRUE)
test = subset(movieData3, sample == FALSE)
```

```
#using Kmeans clustering
```

```
paste("The dimensions of the training data are")
```

```
## [1] "The dimensions of the training data are"
```

```
dim(train)
```

```
## [1] 2249 48
```

```
paste("The dimensions of the testing data are")
```

```
## [1] "The dimensions of the testing data are"
```

```
dim(test)
```

```
## [1] 749 48
```

```
k.means.fit <- kmeans(train[,c(7,12:14,16:26,38)], 20)
k.means.fit$centers
```

```
##      cast_total_facebook_likes      gross      budget      profits
## 1      16549.536 67712868.0 24760156 4.295271e+07
## 2      1173.000 2201412.0 1221550000 -1.221330e+10
## 3      22365.793 263918931.7 183551724 8.036721e+07
## 4        28.000 439162.0 1100000000 -1.099561e+09
## 5     13438.308 27997612.7 74543684 -4.654607e+07
## 6     13191.973 46733581.1 47512568 -7.789872e+05
## 7     16859.471 89263989.2 83361345 5.902645e+06
## 8     10226.813 12394205.4 32164579 -1.977037e+07
## 9     23393.694 62875520.1 155547222 -9.267170e+07
## 10    19563.950 251781661.9 63737500 1.880442e+08
## 11    13626.344 114683644.3 42912295 7.177135e+07
## 12    13164.065 142613315.6 159706522 -1.709321e+07
## 13    15852.517 161229288.7 34210014 1.270193e+08
## 14     6913.298 6061600.6 8344026 -2.282426e+06
## 15    24615.283 181184689.7 115981132 6.520356e+07
## 16     9870.771 33863374.1 16105434 1.775794e+07
## 17    29787.556 364507148.3 133214815 2.312923e+08
## 18       88.000 7292776.5 395000000 -3.877072e+08
## 19    48618.500 641871762.5 202000000 4.398718e+08
## 20    1209.333 901915.3 3033333333 -3.032431e+09
##      num_critic_for_reviews num_user_for_reviews tomatoUserRating
## 1      191.2552      392.9479      3.433854
## 2      363.0000      279.0000      3.200000
## 3      444.8276     1385.5517      3.637931
## 4      150.0000      430.0000      4.000000
## 5      165.2105      314.9850      3.063158
## 6      186.1421      374.8743      3.232240
## 7      217.5546      431.3109      3.279832
## 8      145.7729      208.2125      3.110256
## 9      294.1389      513.9167      3.172222
## 10     243.2500      771.1750      3.642500
## 11     205.3279      450.0984      3.454098
## 12     325.1087      853.5217      3.441304
## 13     227.4000      689.2500      3.670000
## 14     130.7601      187.0746      3.365154
## 15     308.6226      641.9057      3.588679
## 16     170.5229      345.1961      3.349673
## 17     409.7778     1472.2963      3.929630
## 18       71.0000      218.5000      3.250000
## 19     622.1667     2497.1667      4.083333
## 20     149.6667      248.6667      3.833333
##      tomatoRating tomatoReviews tomatoFresh tomatoRotten tomatoUserMeter
## 1      5.906250     123.80729     74.55208     49.25521     67.53646
## 2      3.600000     119.00000     10.00000    109.00000     50.00000
## 3      6.872414     252.44828    184.51724     67.93103     71.03448
```

```
## 4      7.500000      46.00000      40.00000      6.00000      90.00000
## 5      5.097744     126.51128     54.12782     72.38346     50.94737
## 6      5.331148     130.74863     61.02186     69.72678     57.95082
## 7      5.589076     153.18487     80.52941     72.65546     58.42857
## 8      5.067033     109.99634     49.61905     60.37729     51.87179
## 9      5.261111     183.72222     78.58333     105.13889     50.66667
## 10     6.710000     147.92500     102.40000     45.52500     75.85000
## 11     6.291803     140.33607     95.61475     44.72131     69.39344
## 12     6.060870     197.56522     118.26087     79.30435     62.10870
## 13     6.988333     143.20000     109.71667     33.48333     79.63333
## 14     5.926580      92.53160     59.08266     33.44895     63.62075
## 15     6.098113     187.00000     118.39623     68.60377     66.92453
## 16     5.829739     117.78431     72.48693     45.29739     64.96078
## 17     7.192593     243.25926     192.14815     51.11111     80.18519
## 18     5.050000      66.50000     26.50000     40.00000     59.50000
## 19     7.800000     286.66667     249.00000     37.66667     84.16667
## 20     7.633333      84.33333     72.66667     11.66667     87.66667
##      tomatoUserReviews num_voted_users imdbVotes duration
## 1           921602.95         131468.71 134443.70 108.4896
## 2           53348.00          68883.00  93664.00 110.0000
## 3          7407187.72         403357.83 412243.00 136.6552
## 4          147140.00         106160.00 108399.00 124.0000
## 5          149492.98          80849.70  83410.17 119.2030
## 6          353784.74         102738.87 104758.60 113.4481
## 7          538440.55         142402.88 145209.62 115.9496
## 8           76536.48          44776.53  45603.40 109.2234
## 9          200362.17         124343.94 136598.19 120.8611
## 10         3178748.10         367515.78 373808.55 114.4500
## 11         1730981.29         199206.19 203400.23 114.7049
## 12         404829.09         224382.48 240450.39 122.4130
## 13         2142255.23         322003.85 327945.15 120.0167
## 14          57296.52          46001.20  46892.35 102.5057
## 15         2218239.11         276171.17 282376.21 117.6038
## 16         282438.60         101484.71 103488.17 105.6438
## 17         5720295.93         574576.15 589836.19 131.7778
## 18          24083.50          28777.50  28952.50 229.0000
## 19         6913243.67         960746.00 980335.50 165.6667
## 20          98405.33          93554.33  96158.33 126.6667
```

```
# k.means.fit$cluster
k.means.fit$size
```

```
## [1] 192  1 29  1 133 183 119 273  36  40 122  46  60 617  53 306  27
## [18]  2  6  3
```

```
kmeansPrediction <- as.data.frame(k.means.fit$centers)
kmeansPrediction <- kmeansPrediction[,5]
kmeansPrediction
```

```
## [1] 191.2552 363.0000 444.8276 150.0000 165.2105 186.1421 217.5546
## [8] 145.7729 294.1389 243.2500 205.3279 325.1087 227.4000 130.7601
## [15] 308.6226 170.5229 409.7778  71.0000 622.1667 149.6667
```

```

closest.cluster <- function(x) {
  cluster.dist <- apply(k.means.fit$centers, 1, function(y) sqrt(sum((x-y)^2)))
  # print(c( "cluster: " , (which.min(cluster.dist)[1]), kmeansPrediction[which.min(cluster.dist)[1]]))
  # print( kmeansPrediction[which.min(cluster.dist)[1]])
  return (kmeansPrediction[which.min(cluster.dist)[1]])
}

```

```

clusters2 <- apply(test[,c(7,12:14,16:26,38)], 1, closest.cluster)

```

```

# analysis of k means clustering approach

```

```

RSFE_v1 <- test$imdb_score - clusters2

```

```

# RSFE_v1

```

```

RSFE1 <- sum(RSFE_v1)

```

```

# RSFE1

```

```

absRSFE1 <- abs(RSFE1)

```

```

absRSFE1

```

```

## [1] 129724.3

```

```

length(RSFE_v1)

```

```

## [1] 749

```

```

MSFE1 <- absRSFE1 / length(RSFE_v1)

```

```

MSFE1

```

```

## [1] 173.1967

```

```

# calculating the mad

```

```

madtest <- mad(test$imdb_score, center = median(test$imdb_score), constant = 1)

```

```

madKmeans <- mad(clusters2, center = median(clusters2), constant = 1)

```

```

#Using random forests for predicting the IMDB score

```

```

set.seed(7)

```

```

rfdi <- movieData3[sample(nrow(movieData3)), ]

```

```

rf.train <- rfdi[1:2200,]

```

```

rf.test <- rfdi[2201:nrow(rfdi), ]

```

```

paste("The dimensions of the training data are")

```

```

## [1] "The dimensions of the training data are"

```

```

dim(rf.train)

```

```

## [1] 2200 48

```

```

paste("The dimensions of the testing data are")

```

```

## [1] "The dimensions of the testing data are"

```

```
dim(rf.test)
```

```
## [1] 798 48
```

```
set.seed(5)
```

```
rf.rfModel <- randomForest(rfdf$imdb_score ~ rfdf$cast_total_facebook_likes + rfdf$gross + rfdf$budget + rfdf$opening_weekend_box_office)
rf.rfModel
```

```
##
```

```
## Call:
```

```
## randomForest(formula = rfdf$imdb_score ~ rfdf$cast_total_facebook_likes + rfdf$gross + rfdf$budget + rfdf$opening_weekend_box_office,
```

```
##               Type of random forest: regression
```

```
##               Number of trees: 500
```

```
## No. of variables tried at each split: 5
```

```
##
```

```
##               Mean of squared residuals: 0.1498318
```

```
##               % Var explained: 86.29
```

```
#validating the random forest model
```

```
#RMSE
```

```
rf.predictedValues <- predict(rf.rfModel, rfdf)
```

```
# rf.predictedValues
```

```
RSFE_v2 <- rfdf$imdb_score - rf.predictedValues
```

```
# RSFE_v2
```

```
RSFE2 <- sum(RSFE_v2)
```

```
# RSFE2
```

```
absRSFE2 <- abs(RSFE2)
```

```
absRSFE2
```

```
## [1] 13.66584
```

```
length(RSFE_v1)
```

```
## [1] 749
```

```
MSFE2 <- absRSFE2 / length(RSFE_v2)
```

```
mean(rf.rfModel$mse)
```

```
## [1] 0.1554552
```

```
# calculating median absolute deviation
```

```
madtestrf <- mad(rfdf$imdb_score, center = median(rfdf$imdb_score), constant = 1)
```

```
madpredictedrf <- mad(rf.predictedValues, center = median(rf.predictedValues), constant = 1)
```

```
print(c("MSFE for prediction using random forests" , MSFE2))
```

```
## [1] "MSFE for prediction using random forests"
## [2] "0.00455831776740199"
```

```
print(c("MSFE for prediction using K means clustering" , MSFE1))
```

```
## [1] "MSFE for prediction using K means clustering"
## [2] "173.196683616234"
```

```
print(c("MSFE for prediction using Linear Regression" , MSFE))
```

```
## [1] "MSFE for prediction using Linear Regression"
## [2] "1.09274386292992e-14"
```

```
print(c("MAD for Original data:" , madtestrf))
```

```
## [1] "MAD for Original data:" "0.699999999999999"
```

```
print(c("MAD for prediction using random forests" , madpredictedrf))
```

```
## [1] "MAD for prediction using random forests"
## [2] "0.643714999999992"
```

```
print(c("MAD for Original data:" , madtestrf))
```

```
## [1] "MAD for Original data:" "0.699999999999999"
```

```
print(c("MAD for prediction using K means clustering" , madKmeans))
```

```
## [1] "MAD for prediction using K means clustering"
## [2] "24.7499820440997"
```

```
print(c("MAD for Original data:" , Madoriginal))
```

```
## [1] "MAD for Original data:" "0.699999999999999"
```

```
print(c("MAD for prediction using Linear Regression" , MadRegression))
```

```
## [1] "MAD for prediction using Linear Regression"
## [2] "0.680744307507511"
```

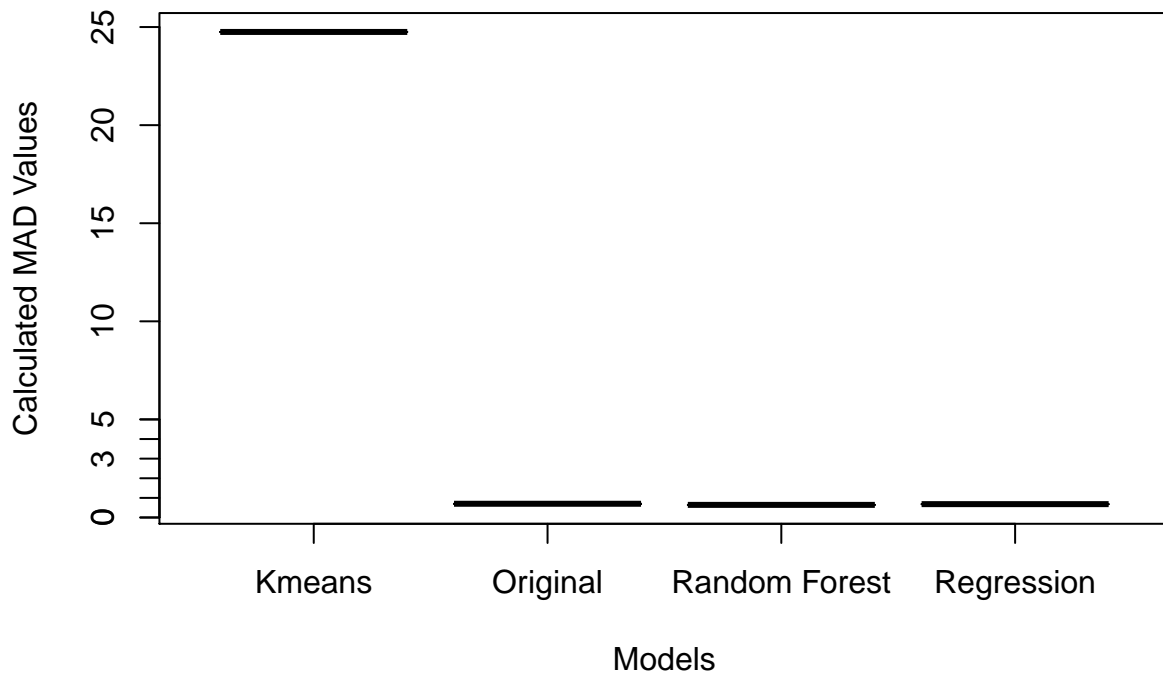
```
# Visualization of results
visualization <- as.data.frame(c(1:6))
visualization$new <- c(1:6)
colnames(visualization) <- c("MSFE","MAD")
row.names(visualization) <- c("Original","Regression","Kmeans","Random Forest","x","y")
visualization[1,1] <- 0
visualization[2,1] <- MSFE
```

```

visualization[3,1] <- MSFE1
visualization[4,1] <- MSFE2
visualization[1,2] <- Madoriginal
visualization[2,2] <- MadRegression
visualization[3,2] <- madKmeans
visualization[4,2] <- madpredictedrf
visualization <- visualization[c(1:4),]
plot(as.factor(rownames(visualization)),visualization$MAD, xlab="Models", ylab="Calculated MAD Values",
axis(2,at = c(0:5))

```

original values vs Observed Values (MAD)

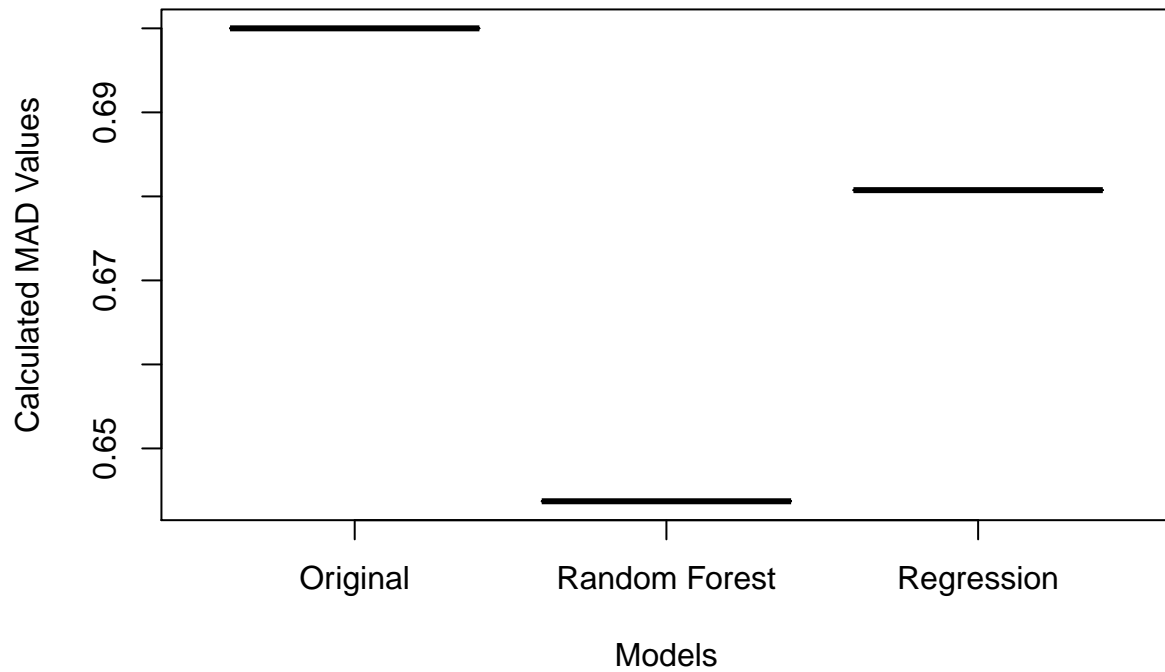


```

plot(as.factor(rownames(visualization[c(1,2,4),])),visualization$MAD[c(1,2,4)], xlab="Models", ylab="Ca

```

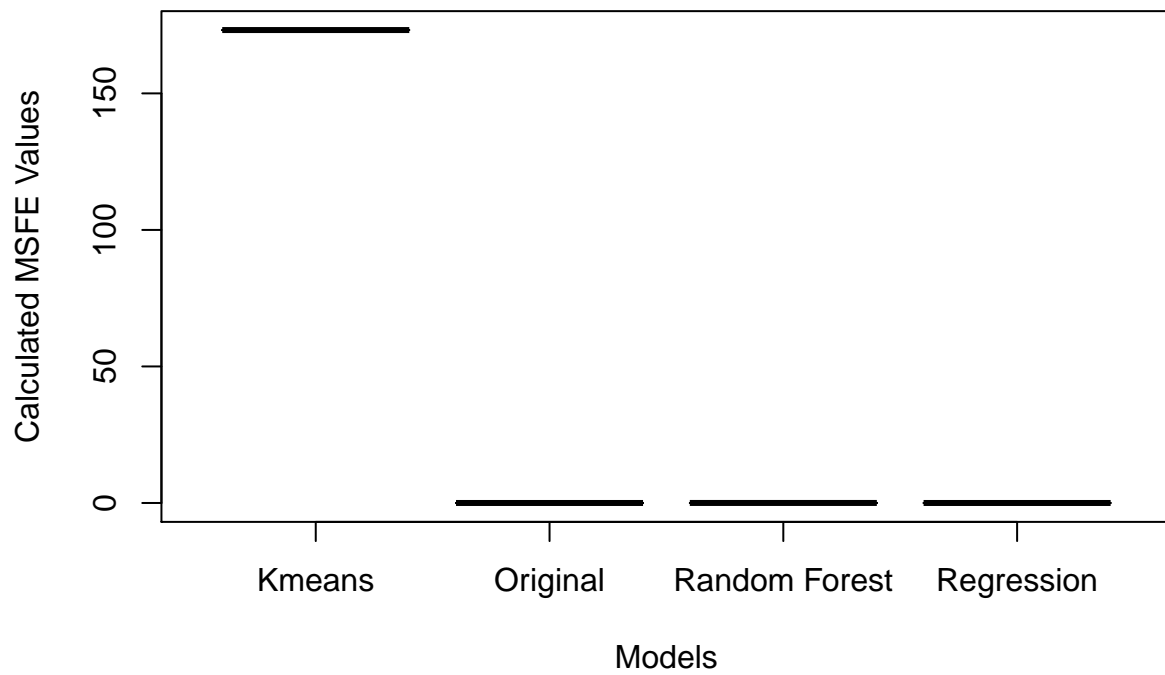
original values vs Observed Values without Kmeans (MAD)




```
# Visualizing MSFE values
```

```
plot(as.factor(rownames(visualization)),visualization$MSFE, xlab="Models", ylab="Calculated MSFE Values")
```

original values vs Observed Values (MSFE)




```
plot(as.factor(rownames(visualization[c(1,2,4),])),visualization$MSFE[c(1,2,4)], xlab="Models", ylab="C
```



```
paste("Looking at the MSFE values for all the three models, we can clearly see that the k-means cluster
```

```
## [1] "Looking at the MSFE values for all the three models, we can clearly see that the k-means cluster
```

```
paste("order of performances: Linear Regression BETTER THAN Random Forests BETTER THAN K Means Clusteri
```

```
## [1] "order of performances: Linear Regression BETTER THAN Random Forests BETTER THAN K Means Cluster
```